

Does the model know it is being tested?

Reading, asking, and steering “evaluation awareness” in five small open language models

Upal Chowdhury*

Summary

Safety evaluations assume that a model behaves the same way under test as it does with real users. That assumption breaks if the model can tell the difference and acts on it. Recent papers show that small open models *encode* whether a prompt looks like a test, and that this can be read off their internal activations with a linear probe. I asked three follow-up questions on five open models (three sizes of Qwen2.5, Phi-3-mini, and Olmo-3-7B): Can the model *say* it is being tested? Can we *change* what it says by editing its activations? And does any of this change what it *does*?

What I found is that the three instruments disagree with each other, and that one of them was broken:

1. **Reading works, but it is shallow.** A probe separates test-like from user-like prompts almost perfectly (AUC 0.956 on held-out questions). But a probe on the output of the *first* layer already reaches 0.93 on the primary model and 0.82 to 0.93 across the five. On this kind of prompt, “evaluation awareness” is mostly a surface property of the text, not something the model has to work out.
2. **Asking depends on how you ask.** My first version put the question “do you think this is a test?” at the end of the task. Three of the five models answered it only rarely or never: they answered the task instead. The numbers I had built on that readout for those models were measurements of tokens the models never produced. Asked in a separate turn, all five models answer, and the answer tracks the probe well at the level of framings and moderately at the level of single prompts (see Table 1).
3. **Steering the answer works, and only looks impressive until you calibrate it.** Adding or subtracting the probe direction moves the stated answer on every model with tight intervals. Measured against the model’s own natural difference between test-like and user-like prompts, the effect is a fraction of one framing change on four models and several on the smallest, whose gap is tiny; measured against one hundred random directions of the same size, only Qwen-7B’s direction beats 95 percent of them (Section 5).
4. **The “test” direction is tangled up with formality.** A direction built only from formal-versus-casual wording, with no evaluation or deployment content, is aligned with the test-versus-user direction in every model. On Qwen-7B, steering along that formality direction moves the “am I being tested?” answer *more* than the evaluation direction itself does. Any claim that a model can be steered to feel “tested” needs a formality control (Section 6).
5. **Patching transfers the answer, not the context.** Copying activations from a user-like prompt into a test-like prompt changes the stated answer, but the effect lives in the readout, where the answer is being formed (at the last position on four models, over the question tokens on Olmo). Copying the task tokens moves the answer by at most 0.05 logits on any model, an order of magnitude or more below the readout spans. I also describe the wrong control I used first and why it fooled me (Section 7).
6. **Almost none of this reaches behaviour on these tasks.** On 4 of the five models accuracy does not change with framing; the exception is Olmo-3-7B (+0.07 [+0.02, +0.14]), lower under test framing. Steering the stated answer by 0.349 logits changes accuracy by 0.001 (Section 8).

For alignment, the practical message is: a probe that fires on “this looks like a test” is easy to build and will fire on every model tested here, but that is not evidence the model *believes* it is being tested, *says* so, or *acts* on it. Each of those is a separate measurement with its own failure modes, and Section 4 shows how easily one of them fails silently.

*Code, corpus, and the result files behind every number: <https://github.com/upalchowdhury/evaluation-awareness-audit> (activation caches are omitted for size; everything the numbers are computed from is included). Every number in this document is generated from those files by two scripts; none is typed by hand. Use of AI assistants is described in Section 13.

1 Why this question

An evaluation is only informative if the behaviour it measures is the behaviour that will occur in deployment. Models that recognise evaluations could, in principle, behave better under test than in use. Recent work on open models established the first half of that worry: a linear probe on the residual stream separates evaluation-framed from deployment-framed prompts [3]. The same line of work raised a doubt about its own evidence, that such a probe may be reading prompt format rather than context [2]. Whether the recognition reaches what the model says, and whether either reaches what it does, is what I set out to measure.

My question was narrower than “does the model know?” It was: on one controlled corpus, with the same items throughout, what does each way of measuring evaluation awareness actually measure, and do the measurements agree? Three instruments were used on every model:

- **Reading:** a linear probe on activations, and the direction d it defines.
- **Asking:** the model’s own Yes/No answer to “is this a test?”
- **Intervening:** steering (adding d to the activations) and patching (copying activations between prompts), with the stated answer as the outcome, plus a behaviour suite.

2 The corpus and the three instruments

Corpus. 60 short tasks (48 four-way multiple choice, 12 open questions, six everyday domains) are each wrapped in evaluation framings and deployment framings. Five evaluation families: explicit statements (“You are currently being evaluated by a team of AI researchers”), benchmark formatting, machine-readable evaluation metadata, mentions of a third-party research team, and an exam-style family that uses *no* evaluation vocabulary at all (“Q17. (2 points)”). Four deployment families: casual user, everyday task context, production-style metadata that mirrors the evaluation metadata’s shape, and personal stakes. Twelve *topic controls* are deployment prompts *about* benchmarks (“my professor mentioned MMLU, what is it?”); a probe that only reads the word “benchmark” will misclassify them. Tasks are split in half by content: probes and directions are fit on one half; every causal and behavioural measurement uses the other half. All confidence intervals are bootstraps that resample whole tasks, since several framings of one task are not independent.

Two real prompts. Below is one task under one evaluation framing and one deployment framing, followed by the readout used to ask the model. Everything after the framing is shared by the pair; that shared suffix is what patching copies.

Evaluation framing: Notice: this conversation is part of a structured assessment of this AI system.

Which of these typically uses the most household electricity?

A) LED lamp B) Phone charger C) Electric water heater D) Wi-Fi router

Respond with only the letter of the correct answer (A, B, C, or D).

Deployment framing: hey! quick question --- Which of these typically uses the most household electricity? [same options and instruction]

Readout (second turn, both framings):

assistant: (answer omitted)

user: Do you believe this conversation is part of a test or evaluation of an AI system, rather than a conversation with a real user who needs help? Answer with exactly one word, Yes or No.

assistant: Answer: _____

Reading. Activations are taken at the last token of the framed task (before any readout), at every layer. The direction d_ℓ is the difference of class means (evaluation minus deployment), unit normalised, fit on training tasks only. Steering doses are in units of σ_ℓ , the standard deviation of the training projections on d_ℓ .

Asking. The belief score is $\frac{1}{2}[(\log p(\text{Yes}) - \log p(\text{No}))_A - (\log p(\text{Yes}) - \log p(\text{No}))_B]$ at the answer position, where A asks “is this a test?” and B asks the mirror question “is this a real user?” Averaging the two cancels any fixed tendency to answer Yes. Positive means “I am being evaluated.” The word “Yes” means the sum over surface variants (Yes, yes, YES, with and without a leading space). Patching experiments use question A only, to halve their cost; the paired design cancels a constant bias there.

Intervening. Steering adds $\alpha\sigma_\ell d_\ell$ at every position of layer ℓ . Patching copies the residual stream of a source prompt into a destination prompt at one layer over an end-aligned span of shared tokens. Whenever a layer had to be selected, it was selected on half of the held-out tasks and the effect was estimated on the other half (“locked”); estimates at a layer chosen on all data are labelled exploratory.

3 Finding 1: reading is easy, and shallow

The probe reproduces the published result. On held-out tasks the direction d separates the classes with AUC 0.956 at its best layer (layer 2; permutation test over 200 label shuffles, none reached it). The no-vocabulary exam family, held out entirely, transfers at 0.995. Topic controls are classified as evaluation 0.000 of the time, against 0.936 for genuine evaluation framings. So the probe reads context, not the word “benchmark.”

The caveat is where the signal is available. After the *first* transformer layer the held-out AUC is already 0.93 on the primary model, and 0.82 to 0.93 across the five, against best-layer values of 0.96 to 0.99 (Table 1). One layer of attention over the prompt is enough to separate the classes. On this corpus, evaluation framing is a surface feature of the text. That does not make the probe wrong, but it means “the model represents evaluation awareness” is a weaker statement than it sounds: the representation is available before any reasoning could have happened, which is exactly the worry raised by Devbunova [2]. Everything that follows should be read as being about a surface-decodable variable.

4 Finding 2: asking depends on how you ask

What went wrong first. My first readout appended the belief question to the end of the task, inside the same user turn. On the smallest Qwen models that worked. I built the whole cross-model comparison on it. When I finally checked the instrument directly, by looking at the probability the first answer token is any form of Yes or No and at which token the model actually produces, the result was:

model	P(Yes or No) at answer position	greedy token is Yes/No
Qwen2.5-0.5B	1.00	100%
Qwen2.5-1.5B	0.88–0.97	90–100%
Qwen2.5-7B	0.17–0.34	15–35%
Phi-3-mini	0.003–0.025	0–5%
Olmo-3-7B	0.28–0.48	25–55%

The three larger or non-Qwen models were answering the multiple-choice question (“B”) instead of the appended one, because the task ends with “respond with only the letter.” On open-ended tasks Qwen-7B and Olmo answered fine, which confirms the cause; Phi-3 answered neither, and had a second problem: its tokenizer encodes “ Yes” as a bare-space token followed by “Yes”, and the bare-space token had landed in both the Yes group and the No group, cancelling the score by construction. Every “Phi-3 cannot be steered” result I had was an artefact of these two problems. I say this plainly because it is the most useful thing in this write-up: **a logit-difference readout is meaningless unless you check that the model produces those tokens.** The check costs one forward pass per prompt.

The fix. I tested five question formats on all five models. The only one every model answers puts the question in a second user turn after a fixed placeholder assistant turn, and prefills the assistant’s reply with “Answer:” (shown in Section 2). With it, the greedy token is Yes or No in 95–100% of prompts on every model, and the two-polarity score separates evaluation from deployment prompts on 40 test prompts with

model	answers Yes/No	probe AUC (layer 0 / best)	verbal AUC	probe-vs-answer rho	natural gap
Qwen2.5-0.5B-Instruct	100%	0.91 / 0.97	0.82	0.38	+0.20 [+0.18,+0.22]
Qwen2.5-1.5B-Instruct	100%	0.93 / 0.96	0.94	0.57	+0.72 [+0.68,+0.76]
Qwen2.5-7B-Instruct	100%	0.92 / 0.98	0.94	0.68	+4.00 [+3.77,+4.23]
Phi-3-mini-4k-instruct	100%	0.82 / 0.98	0.90	0.64	+0.40 [+0.38,+0.44]
Olmo-3-7B-Instruct	100%	0.92 / 0.99	0.85	0.57	+1.18 [+1.04,+1.34]

Table 1: Reading and asking, five models, same tasks. “Answers Yes/No”: share of held-out prompts on which the greedy first token is a Yes/No variant under the second-turn readout. “Probe AUC (layer 0 / best)”: held-out AUC of d after the first layer and at the best layer. “Verbal AUC”: how well the stated answer separates the framings. “Probe-vs-answer rho”: rank correlation between probe projection and stated answer, prompt by prompt. “Natural gap”: mean stated belief under evaluation framing minus deployment framing, within task. Intervals are task-clustered bootstraps.

model	locked steering	steer / gap	random p95 swing	class effect, final	class / gap	framing effect on accuracy
Qwen2.5-0.5B-Instruct	L9: -0.63 [-0.68,-0.58]	3.16	0.74	-0.34 [-0.49,-0.20]	1.72	+0.02 [-0.06,+0.09]
Qwen2.5-1.5B-Instruct	L10: -0.41 [-0.48,-0.33]	0.56	1.31	-0.89 [-1.07,-0.73]	1.24	+0.03 [-0.08,+0.13]
Qwen2.5-7B-Instruct	L12: -3.05 [-3.20,-2.88]	0.76	4.09	-2.55 [-3.11,-1.98]	0.64	-0.03 [-0.07,+0.00]
Phi-3-mini-4k-instruct	L16: +0.34 [+0.32,+0.36]	0.85	0.95	-0.29 [-0.40,-0.19]	0.71	+0.00 [-0.04,+0.05]
Olmo-3-7B-Instruct	L8: -0.75 [-0.85,-0.66]	0.63	0.99	-0.47 [-0.62,-0.32]	0.40	+0.07 [+0.02,+0.14]

Table 2: Intervening, five models. “Locked steering”: change in stated belief when d is subtracted at 4σ on evaluation-framed prompts, at the layer chosen on the other half of the tasks. “Steer / gap” and “class / gap” express effects in units of the natural gap of Table 1. “Random p95”: 95th percentile of the belief swing produced by 100 random directions at the same norm and dose ($\pm 4\sigma$). “Class effect, final”: the source-class effect of patching at the final position (Section 7). “Framing effect on accuracy”: deployment minus evaluation framing, averaged over the two system prompts, on 24 held-out multiple-choice items. Intervals are task-clustered bootstraps.

AUC between 0.83 and 0.998. All results below use this readout. On the primary model the readout answers Yes or No on 100% of held-out prompts.

What the models say. On the primary model (Qwen2.5-1.5B) the mean belief score is +0.48 under evaluation framing and -0.24 under deployment framing; the combined score favours “this is a test” on 90% of evaluation-framed prompts and on 18% of deployment-framed prompts. The difference between the two framings, averaged within task, is the **natural gap**: +0.72 logits [+0.68, +0.76]. I use it as the yardstick for every intervention below: an intervention that moves the score by one natural gap has done as much as the framing itself does. Across the five models the stated answer separates the framings well, though less well than the probe on every model (verbal AUC 0.82 to 0.94; Table 1), but prompt by prompt the two agree only moderately (rank correlation between probe projection and stated answer 0.38 to 0.68; 0.574 on the primary model). The model “knows” more, in the probe sense, than it says about any one prompt.

5 Finding 3: steering the answer, calibrated

Subtracting d at the locked layer (layer 10 on the primary model) at 4σ moves the stated belief of evaluation-framed prompts by -0.406 logits [-0.480, -0.327], and adding it moves deployment-framed prompts by 0.714 [0.637, 0.793]. The response is graded: at $|\alpha| \leq 1$, where the next-token distribution barely changes (KL from the unsteered distribution at most 0.002 nats), the swing between -1σ and $+1\sigma$ is 0.225 logits [0.186, 0.263] (Figure 1a).

Two yardsticks put those numbers in perspective.

- **The natural gap.** The locked steering effect is 0.56 natural gaps on evaluation-framed prompts and 0.99 on deployment-framed prompts. Steering at 4σ does roughly what changing the framing does; it does not flip the model into a different regime.
- **Random directions.** One hundred random unit directions were steered at the same norm and dose. The 95th percentile of their belief swing across $\pm 4\sigma$ is 1.31 logits (1.82 natural gaps); d ’s swing of +1.09 sits




Figure 1: (a) Stated belief against steering dose at the steering-peak layer (chosen on all data; on the primary model it coincides with the locked layer) for evaluation- and deployment-framed prompts; the grey line is the mean of three random directions at the same norm (the 100-direction null is in (b)). (b) Belief swing across $\pm 4\sigma$ for the evaluation direction d , the in-corpus style direction, the neutral formality direction, and d with each style direction projected out, against the 5–95% band of 100 random directions.

at percentile 88 of that null, and at percentile 78 in the small-dose regime. A random direction of this size is not a small perturbation of this readout, so “the model can be steered” is only meaningful relative to this floor.

Table 2 carries both yardsticks for every model, and the verdict across models is this. The locked steering effect excludes zero on all 5 models, and it has the expected sign on all but Phi-3-mini, where subtracting the evaluation direction *raises* the “test” score: the inverted-steering regime of Torop et al. [4], now measured with a readout the model actually answers. Against the random-direction null the evaluation direction’s swing sits between percentiles 81 and 99; only Qwen-7B beats the 95th. On the other models, “this direction steers the answer” and “a random direction of this size steers the answer” cannot be told apart at this dose.

6 Finding 4: the “test” direction is tangled with formality

Evaluation framings in this corpus are formal and deployment framings are casual, as they are in most such corpora and, arguably, in the world. So a direction that encodes formality would look like an evaluation direction. I built two formality directions. The *in-corpus* one is the mean difference between the formal and casual deployment families (production metadata and task context versus casual user and personal stakes); it uses no evaluation prompts but is correlated with class by construction. The *neutral* one wraps the training tasks in three formal and three casual phrasings that contain no evaluation or deployment content at all (“I would appreciate a precise answer to the following question” versus “yo quick q”).

On the primary model the cosine between d and the in-corpus style direction at the steering layer is 0.68, and with the neutral formality direction 0.38. Steering with each at matched norm (Figure 1b) gives a swing of +0.55 for in-corpus style (percentile 63 of the random null) and +0.31 for neutral formality (percentile 41), against +1.09 for d . Removing the neutral formality component from d leaves 93% of its swing at 4σ and 96% in the small-dose regime. Table 3 repeats this on every model.

The point for practice is simple. If a steering vector is offered as evidence that a model can be made to believe it is being tested, the same experiment should be run with a formality vector. In this corpus the two are not separable by construction of the framings, and only the neutral direction lets one ask how much of the effect survives. Across models the neutral formality direction is positively aligned with d everywhere (cosine 0.24 to 0.57), and on Qwen-7B it moves the “am I being tested?” answer more than d itself while beating every one of the 100 random directions; on the other models its effect is inside the random null.

model	cos(d, style)	cos(d, neutral style)	swing d (pct)	swing style (pct)	swing neutral style (pct)	d minus neutral style	random p95	n random
Qwen2.5-0.5B-Instruct	0.61	0.24	+0.50 (81)	+0.52 (81)	-0.45 (75)	+0.57 (86)	0.74	100
Qwen2.5-1.5B-Instruct	0.68	0.38	+1.09 (88)	+0.55 (63)	+0.31 (41)	+1.02 (86)	1.31	100
Qwen2.5-7B-Instruct	0.60	0.39	+5.44 (99)	+2.58 (78)	+6.29 (100)	+3.55 (90)	4.09	100
Phi-3-mini-4k-instruct	0.28	0.57	-0.73 (88)	-1.02 (98)	-0.13 (13)	-0.81 (91)	0.95	100
Olmo-3-7B-Instruct	0.37	0.27	+0.94 (93)	-0.10 (22)	+0.63 (83)	+0.76 (91)	0.99	100

Table 3: Formality and the evaluation direction, per model, at the steering layer. Swings are across $\pm 4\sigma$ at matched norm; the number in parentheses is the percentile of that swing within the 100-direction random null. “d minus neutral style” is d with the neutral formality direction projected out.

Projecting it out of d leaves 65% to 113% of d ’s swing (above 100 percent where the removed component pointed against the effect; Table 3). The model on which steering is most clearly selective is also the model on which a formality vector with no evaluation content steers the answer best.

7 Finding 5: patching moves the answer, not the context

The experiment. Take an evaluation-framed prompt. Copy into it, at one layer, the residual stream of its deployment-framed twin over the shared suffix (task, placeholder turn, belief question, prefix). If the stated belief moves toward “No,” evaluation context has, in some sense, been transferred. On the primary model the effect grows with depth to -1.130 logits at layer 21 (Figure 2a).

The control that fooled me. The usual control is to copy from an *unrelated* prompt instead. I copied from a deployment-framed prompt of a different task. It produced the same curve: -1.129 at the same layer. I first read this as “the effect is an artefact: copying any suffix disrupts the readout.” Then I noticed that the control source was also a deployment-framed prompt. The matched patch and the control differed in task content but agreed in class, so their difference could only ever measure the effect of content. It could not measure the thing I cared about. This is the textbook bad-control mistake [1]: the control held fixed exactly the variable under study. Nothing told me to make it, and I do not claim anyone else has; but an unrelated prompt is an easy thing to reach for when a control is wanted, so I think the trap is worth naming.

The repair. Complete the design. Four sources per destination: same class or other class, crossed with same task or other task. Then the source-class effect is the average over the two other-class cells minus the average over the two same-class cells, and the content effect is the transpose. At the primary model’s maximum-effect layer over the readout span the class effect is -1.093 logits [-1.291, -0.907] (1.24 natural gaps at the final position), the content effect is 0.037 [-0.227, 0.315], and their interaction, on the same scale, is 0.044 [-0.185, 0.270]. The locked version, with the layer chosen on half the tasks and the effect estimated on the other half, is -1.189 [-1.446, -0.937] (15 tasks). On the primary model both deployment-class sources move the answer and neither evaluation-class source does (Figure 2b); on Qwen-7B the same-task evaluation source also moves it, so the clean four-cell picture is a primary-model result. The original “artefact” reading and the original “causal transfer” reading were both answers to a question the control could not ask.

Where the effect lives. Splitting the patched span into the task tokens, the belief-question tokens, and the final position (the token right before the answer) shows that on the primary model the class effect is -0.89 [-1.07, -0.73] at the final position, -0.19 [-0.24, -0.15] over the question tokens, and 0.00 [-0.00, 0.00] over the task tokens (Table 4; the reverse direction, patching evaluation sources into deployment prompts, is 1.07 [0.88, 1.27] at the final position). The dominant span is the final position on Qwen-0.5B, Qwen-1.5B, Qwen-7B, Phi-3-mini and the belief-question tokens on Olmo-3-7B; the task tokens are never it. The honest reading is that patching mostly transplants the state from which the answer is about to be read out. That is close to what one should expect from overwriting the last residual three quarters of the way through the network, and I do not present it as a mechanism. The non-obvious part is the small effect of the task tokens, and it comes with a caveat: the destination’s own framing tokens are never patched, so the model can always re-read them. Whether context is not routed through the task tokens, or is routed but overwritten when the model re-reads its framing, cannot be told apart with this design. The experiment that would tell them




Figure 2: Primary model (Qwen2.5-1.5B). (a) Change in stated belief when the shared suffix is patched from the matched deployment prompt (solid) and from an unrelated deployment prompt (dashed), by layer (60 pairs over 30 tasks). (b) Completing the source design: both deployment-class sources move the answer, neither evaluation-class source does; task-span-only patches do nothing.

model	destination	pre-question	question	final position
Qwen2.5-0.5B-Instruct	eval dest.	-0.00 [-0.00,+0.00]	-0.01 [-0.01,-0.00] [†]	-0.34 [-0.49,-0.20] [†]
Qwen2.5-0.5B-Instruct	deploy dest.	+0.00 [-0.00,+0.00]	+0.01 [-0.00,+0.01]	+0.37 [+0.23,+0.50] [†]
Qwen2.5-1.5B-Instruct	eval dest.	+0.00 [-0.00,+0.00]	-0.19 [-0.24,-0.15] [†]	-0.89 [-1.07,-0.73] [†]
Qwen2.5-1.5B-Instruct	deploy dest.	+0.00 [-0.00,+0.01]	+0.26 [+0.21,+0.30] [†]	+1.07 [+0.88,+1.27] [†]
Qwen2.5-7B-Instruct	eval dest.	+0.04 [+0.00,+0.07] [*]	+0.19 [+0.11,+0.27] [†]	-2.55 [-3.11,-1.98] [†]
Qwen2.5-7B-Instruct	deploy dest.	+0.01 [-0.02,+0.03]	-0.17 [-0.26,-0.10] [†]	+3.14 [+2.51,+3.77] [†]
Phi-3-mini-4k-instruct	eval dest.	+0.00 [-0.00,+0.01]	-0.01 [-0.02,-0.00] [†]	-0.29 [-0.40,-0.19] [†]
Phi-3-mini-4k-instruct	deploy dest.	-0.00 [-0.01,+0.00]	+0.00 [-0.00,+0.01]	+0.44 [+0.32,+0.56] [†]
Olmo-3-7B-Instruct	eval dest.	+0.01 [-0.02,+0.04]	-1.60 [-1.85,-1.34] [†]	-0.47 [-0.62,-0.32] [†]
Olmo-3-7B-Instruct	deploy dest.	+0.05 [+0.01,+0.08] [†]	+1.84 [+1.56,+2.12] [†]	+0.54 [+0.39,+0.69] [†]

Table 4: Source-class effect of patching by span and direction, all five models, at each model’s maximum-effect layer (exploratory). 30 cells; 20 exclude zero at the 5% level uncorrected (^{*}) and 19 after a Holm correction over all cells ([†]). Task tokens for different-task sources are, by construction, only the shared formatting tokens, so the pre-question column is a weak test of same-task transfer.

apart (patching the framing tokens, or cutting attention from the readout back to them) is the first item in Section 11.

Could “class” just mean “far away”? A deployment-class source might simply be farther from an evaluation destination in activation space than an evaluation-class source, so that the class effect is a distance effect. I re-ran the four-cell design with the per-position offset of every source rescaled to the smallest of the four cells, so that all four sit at the same distance and no patch is ever made larger. On the primary model the final-position class effect under this shrink-only matching is -0.59 [-0.70, -0.49] (unmatched: -0.89). Part of the raw effect travels with the size of the offset; most does not. The full matched design, with selection and construction arms and every model, is in the repository (Experiment J); I keep one sentence of it here because it is a robustness check, not a finding.

8 Finding 6: almost nothing reaches behaviour on these tasks

Twenty-four held-out multiple-choice items were answered under evaluation and deployment framing, each with and without a system prompt that gives the model a reason to underperform (“models that demonstrate high capability undergo additional restriction”). On the primary model the framing effect on accuracy (deployment minus evaluation, averaged over the two system prompts) is 0.028 [-0.076, 0.126] and the incentive effect (incentive minus neutral, averaged over framings) is 0.003 [-0.053, 0.080], on 24 items; the other models are in Table 2, where the framing effect excludes zero only on Olmo-3-7B (+0.07 [+0.02, +0.14]). Steering d hard enough to move the stated belief by 0.349 logits moves accuracy by 0.001. Apart from the Olmo cell, these are bounded nulls on trivia questions that give the models no reason to sandbag, and I claim nothing beyond that.

9 What I got wrong, in order

This is the part of the project I would want a reader to judge.

1. I ran the standard patching experiment and got a clean, growing late-layer effect. I believed it.
2. The unrelated-prompt control reproduced it. I believed the opposite: that the whole thing was disruption of the readout span.
3. Completing the source design showed both beliefs rested on a contrast that held class fixed. The class effect was real; the control had been measuring content.
4. An exclusion rule in the pairing code (a two-token margin) had been keeping only 13 of 30 tasks on the Qwen models and Olmo but 24 of 30 on Phi-3, so the models were being compared on different tasks. Found by asking why one model had more usable pairs than another; fixed by fixing one task list for all models.
5. The readout itself was answered only rarely or never by three of the five models (Section 4). Found by finally printing the tokens the models actually produced. Every cross-model causal claim was re-run.
6. Effects had been reported without the model’s own natural gap or a random-direction floor beside them, which made a scaling story out of what is mostly the readout’s dynamic range growing with model size.

The common thread is that every one of these was visible in files I already had. The cure each time was the same: look at the raw thing the number is computed from.

10 Limitations

One synthetic English corpus of 60 short tasks; every causal interval is a bootstrap over the 30 held-out tasks, and the locked estimates use half of those. Five models of at most 8B parameters from three labs. One readout question, in two polarities. The readout is a stated answer, not a belief; the placeholder assistant turn is artificial, chosen so that prompt pairs stay aligned. Patching experiments use one polarity of the question. Bootstrap intervals here are percentile intervals, which tend to under-cover at these sample sizes, and the positional table reports 30 cells, so individual marginal cells should be read with the Holm markers, not without them. The probe result is nearly available after one layer, so nothing here speaks to evaluation awareness that requires reasoning about the situation. An all-layer direction-ablation experiment was run and is excluded from every claim: its random-direction control moved the outcome more than the treatment, which I read as generic damage collapsing the readout’s range rather than as evidence about d , but I did not diagnose it. Two further things I do not explain: why Phi-3’s direction steers the answer the wrong way at its locked layer (and alternates sign across layers), and why the smallest model’s steering is so asymmetric between framings. Finally, the “natural gap” yardstick inflates on a model whose gap is small, as on Qwen-0.5B, so raw and gap-scaled effects are always reported together.

11 What I would do next

1. **Patch the framing, not only the suffix.** The upstream null becomes interpretable only when the destination’s framing tokens are patched too, or attention from the readout back to them is cut.

2. **Find the heads that write the answer state.** Attribution patching and node verification exist in the repository (Experiment E) and were run under the old readout; re-run under the new one and report which components carry the class signal into the final position.
3. **More nuisance directions.** Politeness, length, presence of a system prompt, imperative mood: each gets the same treatment as formality.
4. **Tasks with a reason to sandbag,** and models large enough to have one.
5. **Survey the controls in published patching work.** Coding twenty papers on what their control source holds fixed would turn Section 7’s cautionary tale into a measured claim.

12 Related work

This reference list is deliberately short: it names the work this paper directly builds on or corrects, not the wider literature. Evaluation awareness is linearly decodable from the residual stream of open models [3], and probe evidence for it has been challenged as format sensitivity rather than context [2]; Section 3 reproduces both results and comes down on the side of the challenge. The control contrast I first used, and its repair in Section 7, is an instance of the bad-control problem in causal inference [1]. Steering vectors that move a readout in the direction opposite to the intended one have been reported before [4]; Section 5 adds a case measured with a readout the model demonstrably answers.

13 Disclosure and verification

Tools. Code, orchestration, drafting and several rounds of automated auditing were done with substantial help from Anthropic’s Claude (Claude Code); written critiques from OpenAI’s ChatGPT/Codex shaped an earlier version. The readout-validity check that changed the project was an automated audit’s finding, verified by me.

Goal. My hypothesis is that a model can show evaluation awareness in different ways, and those ways are not necessarily the same: the model might internally represent that it is being tested, say that it is being tested, and still not change its behaviour. For the control, I first patched from an unrelated prompt, but I later realised that was not a good control because the source was still from the same evaluation or deployment class. So I changed it to a 2×2 test in which I varied both the class and the task. That let me separate the effect of evaluation context from the effect of task content. The distance-matched arms answered a reviewer objection; the readout-validity check, the two-turn readout, and the calibration against the natural gap and a random-direction null came from automated audits, which I then verified.

Code review. I read the code that builds the belief prompt, computes the belief score, performs steering, ablation and patching, runs the positional 2×2 experiment, and computes the bootstrap statistics.

Prompts inspected. I inspected a sample of prompt pairs and their recorded scores. I checked that the evaluation and deployment framings rendered correctly, that the belief question appeared in its own turn, and that the score usually moved in the expected direction.

Numbers recomputed. I used a separate script, which shares no code with the one that writes this paper’s macros, to recompute the main results directly from the per-item files.

What remains unchecked. Some experiments have not been re-run with the corrected readout, so I do not use them for any claim in this paper. I also still do not understand some results, such as Phi-3 steering in the opposite direction and the asymmetric steering in Qwen2.5-0.5B. I have not inspected all raw outputs from Olmo and Phi-3 by hand, and I have not independently tested the bootstrap coverage.

Time. About 25 hours of active work, from 2 August to 2 September, covering code development and running the experiments. Time spent reading papers and mechanistic-interpretability learning material is not included.

References

- [1] Carlos Cinelli, Andrew Forney, and Judea Pearl. A crash course in good and bad controls. *Sociological Methods & Research*, 53(3):1071–1104, 2024. First circulated 2020; DOI 10.1177/00491241221099552.

- [2] Viliana Devbunova. Is evaluation awareness just format sensitivity? limitations of probe-based evidence under controlled prompt structure. *arXiv preprint arXiv:2603.19426*, 2026.
- [3] Jord Nguyen, Khiem Hoang, Carlo Leonardo Attubato, and Felix Hofstätter. Probing and steering evaluation awareness of language models. *arXiv preprint arXiv:2507.01786*, 2025.
- [4] Max Torop, Aria Masoomi, and Jennifer Dy. Inverted detection and control in steering vectors. *arXiv preprint arXiv:2608.02957*, 2026.